



Navigation



Machine Learning Mastery

Making Developers Awesome at Machine Learning

[LinkedIn](#)
[Twitter](#)
[Facebook](#)
[YouTube](#)
[Instagram](#)



Your First Deep Learning Project in Python with Keras Step-By-Step

Click to Take the FREE Deep Learning Crash-Course


Search...

[How to Grid Search Hyperparameters for Deep Learning Models in Python With Keras](#)

by Jason Brownlee on August 9, 2017 in Deep Learning

[Regression Tutorial with the Keras Deep Learning Library](#)

[Last Updated on August 27, 2020](#)



Multi-Class Classification Tutorial with the Keras Deep Learning Library

Keras library provides a way to calculate and report on a suite of standard metrics when training deep learning models.

In addition to offering standard metrics for classification and regression problems, Keras also allows you to define and report on your own custom metrics when training deep learning models. This is particularly useful if you want to keep track of a performance measure that better captures the skill of your model during training.

[How to Save and Load Your Keras Deep Learning Model](#)

In this tutorial, you will discover how to use the built-in metrics and how to define and use your own metrics when training deep learning models in Keras.

After completing this tutorial, you will know:

Loving the Tutorials?


- How Keras metrics work and how you can use them when training your models.
- How to use regression and classification metrics in Keras with worked examples.
- How to define and use your own custom metric in Keras with a worked example.

Kick-start >> SEE WHAT'S INSIDE

new book [Deep Learning With Python](#), including *step-by-step tutorials* and the *Python source code* files for all examples.

Let's get started.

- Update Jan/2020:** Updated API for Keras 2.3 and TensorFlow 2.0.



Metrics and How to Use Custom Metrics for Deep Learning with Keras in Python

Photo by [Indi Samarajiva](#), some rights reserved.

Tutorial Overview

This tutorial is divided into 4 parts; they are:

- Keras Metrics
- Keras Regression Metrics
- Keras Classification Metrics
- Custom Metrics in Keras

Keras Metrics

Keras allows you to list the metrics to monitor during the training of your model.

You can do this by specifying the "metrics" argument and providing a list of function names (or function name aliases) to the `compile()` function on your model.

For example:

```
1 model.compile(..., metrics=['mse'])
```

The specific metrics that you list can be the names of Keras functions (like `mean_squared_error`) or string aliases for those functions (like `'mse'`).

Metric values are recorded at the end of each **epoch** on the training dataset. If a validation dataset is also provided, then the metric recorded is also calculated for the validation dataset.

All metrics are reported in verbose output and in the history object returned from calling the `fit()` function. In both cases, the name of the metric function is used as the key for the metric values. In the case of metrics for the validation dataset, the `"val_"` prefix is added to the key.

Both [loss functions](#) and explicitly defined Keras metrics can be used as training metrics.

Keras Regression Metrics

Below is a list of the metrics that you can use in Keras on regression problems.

- Mean Squared Error:** `mean_squared_error`, MSE or mse
- Mean Absolute Error:** `mean_absolute_error`, MAE, mae
- Mean Absolute Percentage Error:** `mean_absolute_percentage_error`, MAPE, mape
- Cosine Proximity:** `cosine_proximity`, cosine

The example below demonstrates these 4 built-in regression metrics on a simple contrived regression problem.

```
1 from numpy import array
```

Start Machine Learning

1/21

https://machinelearningmastery.com/custom-metrics-deep-learning-keras-python/

```

2 from keras.models import Sequential
3 from keras.layers import Dense
4 from matplotlib import pyplot
5 # prepare sequence
6 X = array([0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0])
7 # create model
8 model = Sequential()
9 model.add(Dense(2, input_dim=1))
10 model.add(Dense(1))
11 model.compile(loss='mse', optimizer='adam', metrics=['mse', 'mae', 'mape', 'cosine'])
12 # train model
13 history = model.fit(X, y, epochs=500, batch_size=len(X), verbose=2)
14 # plot metrics
15 pyplot.plot(history.history['mean_squared_error'])
16 pyplot.plot(history.history['mean_absolute_error'])
17 pyplot.plot(history.history['mean_absolute_percentage_error'])
18 pyplot.plot(history.history['cosine_proximity'])
19 pyplot.show()

```

Note: Your results may vary given the stochastic nature of the algorithm or evaluation procedure, or differences in numerical precision. Consider running the example multiple times.

One Regression Tutorial with the Keras Deep Learning Library in Python

Running the example prints the metric values at the end of each epoch.

```


1 Multi-Class Classification Tutorial with the
2 Epoch 96/100
3 0s - loss: 1.0596e-04 - mean_squared_error: 1.0596e-04 - mean_absolute_error: 0.0088 - mean_absolute_percentage_error: 3.5611 - cosine_proximity: -1.0000e+00
4 Epoch 97/100
5 0s - loss: 1.0354e-04 - mean_squared_error: 1.0354e-04 - mean_absolute_error: 0.0087 - mean_absolute_percentage_error: 3.5178 - cosine_proximity: -1.0000e+00
6 Epoch 98/100
7 0s - loss: 1.0116e-04 - mean_squared_error: 1.0116e-04 - mean_absolute_error: 0.0086 - mean_absolute_percentage_error: 3.4738 - cosine_proximity: -1.0000e+00
8 Epoch 99/100
9 0s - loss: 9.8820e-05 - mean_squared_error: 9.8820e-05 - mean_absolute_error: 0.0085 - mean_absolute_percentage_error: 3.4294 - cosine_proximity: -1.0000e+00
10 Epoch 100/100
11 0s - loss: 9.6515e-05 - mean_squared_error: 9.6515e-05 - mean_absolute_error: 0.0084 - mean_absolute_percentage_error: 3.3847 - cosine_proximity: -1.0000e+00

```

A line plot showing the metrics over training epochs is then created.

The Deep Learning with Python EBook is where you'll find the *Really Good* stuff.

>> SEE WHAT'S INSIDE



Note that the metrics were specified using string alias values ['mse', 'mae', 'mape', 'cosine'] and were referenced as key values on the history object using their expanded function name.

We could also specify the metrics using their expanded name, as follows:

```
1 model.compile(loss='mse', optimizer='adam', metrics=['mean_squared_error', 'mean_absolute_error', 'mean_absolute_percentage_error', 'cosine_proximity'])
```

We can also specify the function names directly if they are imported into the script.

```

1 from keras import metrics
2 model.compile(loss='mse', optimizer='adam', metrics=[metrics.mean_squared_error, metrics.mean_absolute_error, metrics.mean_absolute_percentage_error, metrics.cosine_proximity])

```

You can also use the loss functions as metrics.

For example, you could use the Mean squared Logarithmic Error (*mean_squared_logarithmic_error*, *MSLE* or *msle*) loss function as a metric as follows:

```
1 model.compile(loss='mse', optimizer='adam', metrics=['msle'])
```

Keras Classification Metrics

Below is a list of the metrics that you can use in Keras on classification problems.

- **Binary Accuracy:** binary_accuracy, acc
- **Categorical Accuracy:** categorical_accuracy, acc
- **Sparse Categorical Accuracy:** sparse_categorical_accuracy
- **Top k Categorical Accuracy:** top_k_categorical_accuracy (requires you specify a k parameter)
- **Sparse Top k Categorical Accuracy:** sparse_top_k_categorical_accuracy (requires you specify a k parameter)

Accuracy is special.

Regardless of whether your problem is a binary or multi-class classification problem, you can specify the 'accuracy' metric to report on accuracy.

Below is an example of a binary classification problem with the built-in accuracy metric demonstrated.

```

1 from numpy import array
2 from keras.models import Sequential
3 from keras.layers import Dense
4 from matplotlib import pyplot
5 # prepare sequence
6 X = array([0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0])
7 y = array([0, 0, 0, 0, 0, 1, 1, 1, 1, 1])
8 # create model
9 model = Sequential()
10 model.add(Dense(2, input_dim=1))
11 model.add(Dense(1, activation='sigmoid'))
12 model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
13 # train model
14 history = model.fit(X, y, epochs=400, batch_size=len(X), verbose=2)
15 # plot metrics
16 pyplot.plot(history.history['accuracy'])
17 pyplot.show()

```

Start Machine Learning

You can master applied Machine Learning without math or fancy degrees. Find out how in this *free* and *practical* course.

START MY EMAIL COURSE

Start Machine Learning

Note: Your results may vary given the stochastic nature of the algorithm or evaluation procedure, or differences in numerical precision. Consider running the example a few times and compare the average outcome.



Running the example reports accuracy at the end of each training epoch.

```
1 Epoch 396/400
2 0s - loss: 0.5934 - acc: 0.9000
3 Epoch 397/400
4 0s - loss: 0.5932 - acc: 0.9000
5 Epoch 398/400
6 0s - loss: 0.5930 - acc: 0.9000
7 Epoch 399/400
8 0s - loss: 0.5927 - acc: 0.9000
9 Epoch 400/400
10 0s - loss: 0.5925 - acc: 0.9000
11 Keras: 0.5925 - acc: 0.9000
```

A line plot of accuracy over epoch is created.



[Regression Tutorial with the Keras Deep Learning Library in Python](#)



[Multi-Class Classification Tutorial with the Keras Deep Learning Library](#)




[How to Save and Load Your Keras Deep Learning Model](#)

Loving the Tutorials?

The *Deep Learning with Python* EBook is where you'll find the *Really Good* stuff.

[>> SEE WHAT'S INSIDE](#)



Line Plot of Built-in Keras Metrics for Classification

Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**. Find out how in this *free* and *practical* course.

[START MY EMAIL COURSE](#)

Custom Metrics in Keras

You can also define your own metrics and specify the function name in the list of functions for the *"metrics"* argument when calling the *compile()* function.

A metric I often like to keep track of is Root Mean Square Error, or RMSE.

You can get an idea of how to write a custom metric by examining the code for an existing metric.

For example, below is the code for the [mean_squared_error loss function and metric in Keras](#).

```
1 def mean_squared_error(y_true, y_pred):
2     return K.mean(K.square(y_pred - y_true), axis=-1)
```

K is the backend used by Keras.

From this example and other examples of loss functions and metrics, the approach is to use standard math functions on the backend to calculate the metric of interest.

For example, we can write a custom metric to calculate RMSE as follows:

```
1 from keras import backend
2
3 def rmse(y_true, y_pred):
4     return backend.sqrt(backend.mean(backend.square(y_pred - y_true), axis=-1))
```

You can see the function is the same code as MSE with the addition of the *sqrt()* wrapping the result.

We can test this in our regression example as follows. Note that we simply list the function name directly rather than providing it as a string or alias for Keras to resolve.

```
1 from numpy import array
2 from keras.models import Sequential
3 from keras.layers import Dense
4 from matplotlib import pyplot
5 from keras import backend
6
7 def rmse(y_true, y_pred):
8     return backend.sqrt(backend.mean(backend.square(y_pred - y_true), axis=-1))
9
10 # prepare sequence
11 X = array([0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0])
12 # create model
13 model = Sequential()
14 model.add(Dense(2, input_dim=1, activation='relu'))
15 model.add(Dense(1))
16 model.compile(loss='mse', optimizer='adam', metrics=[rmse])
17 # train model
18 history = model.fit(X, X, epochs=500, batch_size=len(X), verbose=2)
19 # plot metrics
20 pyplot.plot(history.history['rmse'])
21 pyplot.show()
```

Note: Your results may vary given the stochastic nature of the algorithm or evaluation procedure, or differences in numerical precision. Consider running the example a few times and compare the average outcome.

Running the example reports the custom RMSE metric at the end of each training epoch.

```
1 ...
2 Epoch 496/500
3 0s - loss: 1.2992e-06 - rmse: 9.7909e-04
4 Epoch 497/500
5 0s - loss: 1.2681e-06 - rmse: 9.6731e-04
6 Epoch 498/500
7 0s - loss: 1.2377e-06 - rmse: 9.5562e-04
8 Epoch 499/500
9 0s - loss: 1.2079e-06 - rmse: 9.4403e-04
10 Epoch 500/500
11 0s - loss: 1.1788e-06 - rmse: 9.3261e-04
```






At the end of the run, a line plot of the custom RMSE metric is created.


[Start Machine Learning](#)

Never miss a tutorial:



Picked for you:

-  Your First Deep Learning Project in Python with Keras Step-By-Step
-  How to Grid Search Hyperparameters for Deep Learning Models in Python With Keras
-  Regression Tutorial with the Keras Deep Learning Library in Python
-  Multi-Class Classification Tutorial with the Keras Deep Learning Library
-  How to Save and Load Your Keras Deep Learning Model



Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**. Find out how in this *free* and *practical* course.

START MY EMAIL COURSE

Your custom metric function must operate on Keras internal data structures that may be different depending on the backend used (e.g. *tensorflow.python.framework.ops.Tensor* when using tensorflow) rather than the raw *yhat* and *y* values directly.

Loving the Tutorials?

For this reason, I would recommend using the backend math functions wherever possible for consistency and execution speed. The *Deep Learning with Python* EBook is where you'll find the *Really Good* stuff.

Further Reading

>> SEE WHAT'S INSIDE

This section provides more resources on the topic if you are looking go deeper.

- [Keras Metrics API documentation](#)
- [Keras Metrics Source Code](#)
- [Keras Loss API documentation](#)
- [Keras Loss Source Code](#)

Summary

In this tutorial, you discovered how to use Keras metrics when training your deep learning models.

Specifically, you learned:

- How Keras metrics works and how you configure your models to report on metrics during training.
- How to use classification and regression metrics built into Keras.
- How to define and report on your own custom metrics efficiently while training your deep learning models.

Do you have any questions?

Ask your questions in the comments below and I will do my best to answer.

Develop Deep Learning Projects with Python!

What If You Could Develop A Network in Minutes

...with just a few lines of Python


Discover how in my new Ebook:
[Deep Learning With Python](#)

It covers **end-to-end projects** on topics like:
Multilayer Perceptrons, Convolutional Nets and Recurrent Neural Nets, and more...

Finally Bring Deep Learning To Your Own Projects

Skip the Academics. Just Results.

SEE WHAT'S INSIDE



[Tweet](#) [Share](#) [in Share](#)

About Jason Brownlee

Jason Brownlee, PhD is a machine learning specialist who teaches developers how to get results with modern machine learning methods via hands-on tutorials.

[View all posts by Jason Brownlee](#) →

< 10 Command Line Recipes for Deep Learning on Amazon Web Services

Get the Most out of LSTMs on Your Sequence Prediction Problem >

153 Responses to *How to Use Metrics for Deep Learning with Keras in Python*

Gerrit Govaerts August 9, 2017 at 5:03 pm #

Off topic but interesting none the less :

1) how to train an ensemble of models in the same time it takes to train 1
<http://www.kdnuggets.com/2017/08/train-deep-learning-faster-snapshot-ensembling.html>

2) when not to use deep learning
<http://www.kdnuggets.com/2017/07/when-not-use-deep-learning.html>

REPLY ↩


Start Machine Learning

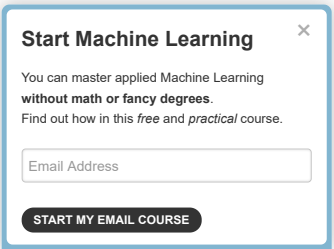
Never miss a tutorial:
Jason Brownlee August 10, 2017 at 6:51 am #



REPLY ↩

Picked for you:

**Oliver** August 14, 2017 at 6:08 pm #
You've just learned how to use Keras with Keras Step-By-Step Hi Jason,
Thanks again for another great topic on keras but I'm a R user !
[How to Grid Search Hyperparameters for Keras](#)
I work with keras on R, but how about to implement custom metric 'rmse' on keras R please ?
Because I find something like that on the github repository :
metric_mean_squared_error <- function(y_true, y_pred) {
 # https://en.wikipedia.org/wiki/Root-mean-squared_error
 Sqrt(mean((y_true - y_pred)^2))
}
attr(metric_mean_squared_error, "py_function_name") <- "mean_squared_error"
I'm poor at my code
[Keras Deep Learning Library](#)
K\$rmse(K\$mean(K\$square(y_pred - y_true)))
[How to Save and Load Your Keras Deep Learning Model](#) is returned)



REPLY ↩

John August 14, 2017 at 10:28 pm #
The [Deep Learning with Python](#) Book is where you'll find the *Really Good* stuff.
Ok finally I make it return a value different from 'nan', but the result is not the same as the square root of 'mse' from keras ?? Maybe due to the arg 'axis = -1' ?
[>> SEE WHAT'S INSIDE](#)

REPLY ↩

Jason Brownlee August 12, 2017 at 6:46 am #
Sorry, I have not used Keras in R, I don't have good advice for you at this stage.

REPLY ↩

John December 14, 2017 at 10:50 pm #
Hi Jason,
Thanks for your very good topic on evaluation metrics in keras. can you please tell me how to compute macro-F and the micro-F scores?
thanks in advance

REPLY ↩

Jason Brownlee December 15, 2017 at 5:33 am #
Sorry, I am not familiar with those scores John.
Perhaps find a definition and code them yourself?

REPLY ↩

Stepan Lavrinenko February 16, 2020 at 4:10 am #
Have you considered using the scikit-learn's implementations? https://scikit-learn.org/stable/modules/generated/sklearn.metrics.f1_score.html
Or would these not work with tensorflow?

REPLY ↩

Jason Brownlee February 16, 2020 at 6:14 am #
See here:
<https://machinelearningmastery.com/how-to-calculate-precision-recall-f1-and-more-for-deep-learning-models/>

REPLY ↩

Linda Cen December 22, 2017 at 5:26 am #
Hi Jason,
I used your "def rmse" in my code, but it returns the same result of mse.

```
# define data and target value
X = TFIDF_Array
Y = df['Shrinkage']

# custom metric to calculate RMSE
def RMSE(y_true, y_pred):
    return backend.sqrt(backend.mean(backend.square(y_pred - y_true), axis=-1))

# define base model
def regression_model():
    # create model
    model = Sequential()
    model.add(Dense(512, input_dim=X.shape[1], kernel_initializer='uniform', activation='relu'))
    model.add(Dropout(0.5))
    model.add(Dense(1, kernel_initializer='uniform'))
    # compile model
    model.compile(loss='mse', optimizer='sgd', metrics=[RMSE])
    return model

# evaluate model
estimator = KerasRegressor(build_fn=regression_model, nb_epoch=100, batch_size=32, verbose=0)
kfold = KFold(n_splits=3, random_state=1)
reg_results = cross_val_score(estimator, X, Y, cv=kfold)
```

REPLY ↩

Jason Brownlee December 22, 2017 at 5:37 am #
[Start Machine Learning](#)

REPLY ↩


Did the example in the post – copied exactly – work for you?
Never miss a tutorial:





Linda Cen December 22, 2017 at 8:58 am #

REPLY ↩

Picked for you:

- 

Your file [3922664 Python with Keras Deep Learning](#) was picked for you.
Epoch 497/500
0s – loss: 3.8870e-04 – rmse: 0.0169
Epoch 498/500
0s – loss: 3.8518e-04 – rmse: 0.0169
Epoch 499/500
0s – loss: 3.8169e-04 – rmse: 0.0168
Epoch 500/500
0s – loss: 3.7821e-04 – rmse: 0.0167
Regression Tutorial with the Keras Deep Learning Library
It gave back different values from yours.
- 

Multi-Class Classification Tutorial with the Keras Deep Learning Library
Linda Cen December 22, 2017 at 9:04 am #
- 

How to Save and Load Your Keras Deep Learning Model
Epoch 497/500
0s – loss: 0.0198 – mean_squared_error: 0.0198
Epoch 498/500
0s – loss: 0.0197 – mean_squared_error: 0.0197
Epoch 499/500
0s – loss: 0.0197 – mean_squared_error: 0.0197
Epoch 500/500
0s – loss: 0.0196 – mean_squared_error: 0.0196
The Deep Learning Library really is **Really Good** stuff.
where you'll find the **Really Good** stuff.
and these were the result when I used:
>> SEE WHAT'S INSIDE [mean_squared_error]
I didn't see any difference of MSE and RMSE here.
Please advise. Thanks.

Jason Brownlee December 22, 2017 at 4:15 pm #

REPLY ↩

Yes, this is to be expected. Machine learning algorithms are stochastic meaning that the same algorithm on the same data will give different results each time it is run. See this post for more details:
<https://machinelearningmastery.com/randomness-in-machine-learning/>

lila January 30, 2018 at 5:19 am #

REPLY ↩

Dear Jason,
Thank you again for the awesome blog and clear explanations
If I understood well, RMSE should be equal to sqrt(mse), but this is not the case for my data:
Epoch 130/1000
10/200 [>.....] – ETA: 0s – loss: 0.0989 – rmse: 0.2656
200/200 [=====] – 0s 64us/step – loss: 0.2856 – rmse: 0.4070
Please sir, how can we calculate the coefficient of determination

Jason Brownlee January 30, 2018 at 9:56 am #

REPLY ↩

The mse may be calculated at the end of each batch, the rmse may be calculated at the end of the epoch because it is a metric.

Josh Park March 27, 2019 at 11:14 am #

REPLY ↩

Hi Jason, thanks for the helpful blog. Quick question regarding your reply here, if the rmse metric is calculated at the end of each epoch, why is it constantly being updated during an epoch whenever you're training?

Jason Brownlee March 27, 2019 at 2:06 pm #

REPLY ↩

It is calculated/estimated per batch I believe.

Josh Park March 27, 2019 at 10:47 pm #

Thanks for your reply. If that's the case, why is the square root of the MSE loss function not equal to the RMSE metric value from above if they are both calculated at the end of each batch?

Jason Brownlee March 28, 2019 at 8:15 am #

They should be, and if not, then there is a difference in the samples used to calculate the score – e.g. batch vs epoch, or a difference in precision between the two calculations causing rounding errors.
You could try digging into the code if this matters.
Generally, I recommend a separate standalone evaluation of model performance and only use training values as a rough/directional assessment.

lila January 30, 2018 at 5:26 am #

REPLY ↩

For the determination coefficient I use this basic code
S1, S2 = 0, 0
for i in range(len(Y)):
S1 = S1 + (Y_pred_array[i] – mean_y)**2
S2 = S2 + (Y_array[i] – mean_y)**2

Start Machine Learning

R2 = S1/S2

Never miss a tutorial:

But this gives give bad results

in

tw

f

en

rs

Picked for you

Valid March 6, 2018 at 6:21 am #

Your First Deep Learning Project in Python with Keras Step-By-Step

How can you deal with y_pred as iterable also it is a Tensor?

Thanks

How to Grid Search Hyperparameters for Deep Learning Models in Python With Keras

Amilkar Herrera January 8, 2021 at 4:06 pm #

Regression Tutorial with the Keras Deep Learning Library

I've been having this same problem. Readin the documentation you can see that by default, the metrics are evaluated by batch and the averaged. In this case, the scalar metric value you are tracking during training and evaluation is the average of the per-batch metric values for all batches see during a given epoch. For the details, see <https://keras.io/api/metrics/>

Multi-Class Classification Tutorial with the Keras Deep Learning Library

Jason Brownlee January 9, 2021 at 6:38 am #

How to Save and Load Your Keras Deep Learning Model

My advice is to calculate the metric manually via the evaluate() function to get a true estimate of model performance. Any scores reporting during training are just a rough approximation.

Start Machine Learning

X

You can master applied Machine Learning without math or fancy degrees. Find out how in this free and practical course.

Email Address

START MY EMAIL COURSE

Loving the Tutorials?

Sam February 28, 2018 at 5:58 pm #

The Deep Learning with Python Book is where you'll find the Really Good stuff. Thanks for the article. How does Keras compute a mean statistic in a per batch fashion? Does it internally (magically) aggregate the sum and count to that point in the epoch and print the measure or does it compute the scalar metric value you are tracking during training and evaluation is the average of the per-batch metric values for all batches see during a given epoch and then again re-compute the metric at the end of each epoch over the entire data?

>> SEE WHAT'S INSIDE

REPLY

Jason Brownlee March 1, 2018 at 6:09 am #

I believe the sum is accumulated and printed at the end of each batch or end of each epoch. I don't recall which.

REPLY

Valid March 6, 2018 at 6:18 am #

Great post and just in time as usual;

The issue is that I am trying to calculate the loss based on IoU (Intersection over union) and I have no clue how to do it using my backend (TensorFlow)

My output is like this(xmin,ymin,xmax,ymax)

Thanks

REPLY

Jason Brownlee March 6, 2018 at 6:20 am #

Sorry, I have not implemented (or heard of) that metric.

REPLY

MLT March 8, 2018 at 8:39 am #

```
model.compile(loss='mse', optimizer='adam', metrics=[rmse])
```

Epoch 496/500
0s - loss: 1.2992e-06 - rmse: 9.7909e-04

loss is mse. Should mse = rmse^2? Above value (9.7909e-04)^2 is 9.6e-8, which mismatch 1.2992e-06. Did I misunderstand something? Thanks.

REPLY

Jason Brownlee March 8, 2018 at 2:54 pm #

The loss and metrics might not be calculated at the same time, e.g. end of batch vs end of epoch.

REPLY

MLT March 9, 2018 at 8:06 am #

Thanks for reply.

history = model.fit(X, X, epochs=500, batch_size=len(X), verbose=2)

I thought the duration of batch is equal to one epoch, since batch_size=len(X). If it is correct?

Furthermore, it seems that the loss of epoch is also updated each iteration.

Epoch 496/500
0s - loss: 1.2992e-06 - rmse: 9.7909e-04

REPLY

Jason Brownlee March 10, 2018 at 6:13 am #

No, one epoch is comprised of 1 or more batches. Often 32 samples per batch are used as a default.

Lear more here:
<https://machinelearningmastery.com/gentle-introduction-mini-batch-gradient-descent-configure-batch-size/>

REPLY

MLT March 10, 2018 at 7:38 am #

Thanks a lot for your time to explain and find the link.

I am sorry. I think I did not express my thoughts correctly.

In the above example, history = model.fit(X, X, epochs=500, batch_size=len(X), verbose=2)

REPLY

Start Machine Learning

^

Never miss a tutorial:

in

tw

f

on

wh


rs

July 24, 2018 at 6:13 am #

REPLY ↩


You will need to work with Tensor types. That is the expectation of Keras.

Picked for you:




[Your First Deep Learning Project in Python with Keras Step-By-Step](#)
Anam August 11, 2018 at 8:04 pm #

REPLY ↩




Dear Jason,
can we use precision and recall metrics for Deep Learning with Keras in Python?
Deep Learning Models in Python With Keras in advance.




[Regression Tutorial with the Keras Deep Learning Library](#)
Jason Brownlee August 12, 2018 at 6:32 am #

REPLY ↩



[Multi-Class Classification Tutorial with the Keras Deep Learning Library](#)
Omar August 30, 2018 at 8:17 pm #

REPLY ↩



[How to Save and Load Your Keras Deep Learning Model](#)
Hello Mr Jason
I have a question that have confused me for so long.
For a multiple output regression problem, what does the MSE loss function compute exactly ?
Thank you in advance.
Loving the Tutorials?
The [Deep Learning with Python](#) EBook is where you'll find the *Really Good* stuff.
Nitin Pasumarthy September 1, 2018 at 3:47 am #

>> SEE WHAT'S INSIDE

```
At run time, I wanted to bucket the classes and evaluate. So tried this function but it returns nan.

def my_metric(y_true, y_pred):
    actual = tf.floor( y_true / 10 )
    predicted = tf.floor( y_pred / 10 )
    return K.categorical_crossentropy(actual, predicted)
```

[Jason Brownlee](#) September 1, 2018 at 6:22 am #

REPLY ↩

Sorry, I cannot debug your code. Perhaps post to stackoverflow?

[Nitin Pasumarthy](#) September 1, 2018 at 9:06 am #

REPLY ↩

Sorry. Thanks a lot. Learned some good things 😊

[Omar](#) September 26, 2018 at 1:56 am #

REPLY ↩

Hello mr Jason
For a multiple output regression problem, what does the MSE loss function compute exactly ?
Is it the sum of the MSE over all the output variables, the average or something else ?
Thank you in advance.

[Jason Brownlee](#) September 26, 2018 at 6:17 am #

REPLY ↩

The average of the squared differences between model predictions and true values.

[Omar](#) October 1, 2018 at 2:36 am #

REPLY ↩

This is for just one output, what if I have multiple outputs ?

[Jason Brownlee](#) October 1, 2018 at 6:27 am #

REPLY ↩

You can calculate the metric for each time step or output. I have an example here:
<https://machinelearningmastery.com/multi-step-time-series-forecasting-long-short-term-memory-networks-python/>

[JG](#) September 30, 2018 at 2:39 am #

REPLY ↩

Thank you so much for your Tutorial. Nowadays I follow your twitter proposals everyday. It's great !
I have two questions:
1) regarding sequential model in the last example;
if I remove activation definition = 'relu', in your last code example ... I got a surprising better RMSE performance values... it is suggest to me that has something to do with Regression issues that works better if we do not put activation at all in the first hide layer. Is it casual result or any profound reason?
2) using a same architectural model, which is better a Regression approach (we leave out the activation in the output layer) or a multinomial classification (we set up the appropriate 'softmax' as activation in the output layer), imagine for example, we analyze same problem, e.g. we have all continuos label output or any discrete multiclass label (for getting for example rounded real number by their equivalent integer number), for a serie of real number samples ...I mean is there any intrinsic advantage or behavior using Regression analysis vs Multinomial classification ?
thanks
JG

[Jason Brownlee](#) September 30, 2018 at 6:07 am #

REPLY ↩

It really depends on the problem as to the choice and benefit of activation functions.

Start Machine Learning

^

https://machinelearningmastery.com/custom-metrics-deep-learning-keras-python/


9/21


In terms of activation in the output layer – what I think you're asking about, the heuristics are:


Never miss a tutorial:


- regression: use 'linear'
- binary classification: use 'sigmoid'
- multi-class classification: use 'softmax'.


Picked for you:

 **Your First Deep Learning Project in Python with Keras Step-By-Step**
JG September 30, 2018 at 8:20 am #

 How to imagine a continuous curve for a linear regression as output of prediction vs imagine for the same problem different outputs of segments to categorize a big multi-class classification, using the same main model structure (obviously with different units and activation at output, loss and metrics at compilation, etc) ...which model will perform better the regression or the multi-classification? My intuition tell me that multi-class it is more fine because it can focus on specific segment output (classes) of the linear regression curve (and even it has more units at output). I try in the future to experiment with some specific examples, to search for my question.

 **Regression Tutorial with the Keras Deep Learning Library in Python**
way, do you like women basket? congratulations Australia won to España (-Spain it is my country -) two hours ago...

 **Multi-Class Classification Tutorial with the Keras Deep Learning Library**
Jason Brownlee October 1, 2018 at 6:23 am #
I don't follow, sorry.

 A classification model is best for classification and will perform beyond poorly for regression.
Learning Model
not a sports guy, sorry 😊

Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**. Find out how in this *free* and *practical* course.

START MY EMAIL COURSE

Loving the Tutorials?

Ronny November 2, 2018 at 4:53 am #

The **Deep Learning with Python** EBook is where you find the **Really Good** stuff.

I follow **>> SEE WHAT'S INSIDE** my own metric function successfully. I was also able to plot it.

The problem that I encountered was when I tried to load the model and the saved weights in order to use `model.evaluate_generator()`. I keep getting the error: "Exception has occurred: ValueError too many values to unpack (expected 2) "

I was wondering if you know how to solve this problem.

Jason Brownlee November 2, 2018 at 5:59 am #

I have not seen this. Is your version of Keras up to date? v2.2.4 or better?

Paul A. Gureghian November 4, 2018 at 9:08 am #

How to extract and store the accuracy output from 'loss' and 'metrics' in the `model.compile` step in order to pass those float values to `mlflow's log_metric()` function ?

Paul A. Gureghian November 4, 2018 at 9:10 am #

```
history = regr.compile(optimizer, loss = 'mean_squared_error', metrics=['mae'])
```

My 'history' variable keeps coming up as 'None type'

Jason Brownlee November 5, 2018 at 6:07 am #

That is odd, I have not seen that before.

Perhaps post to the keras user group:

<https://machinelearningmastery.com/get-help-with-keras/>

Orion November 5, 2018 at 12:46 am #

Hi Dr. Brownlee,

By definition, rmse should be square root of mse.

But if we fit keras with batches, rmse would not be calculated correctly.

Could you give me some advices about how to use customized rmse metric with batches fitting properly?

Jason Brownlee November 5, 2018 at 6:16 am #

If you add RMSE as a metric, it will be calculated at the end of each epoch, i.e. correctly.

Rohan November 8, 2018 at 4:26 am #

Why is the cosine proximity value negative in this case. Should it not be positive since the dot product computed is of the same vectors it should be +1.0 right?

Jason Brownlee November 8, 2018 at 6:13 am #

Perhaps because the framework expects to minimize loss.

Rohan November 8, 2018 at 7:03 am #

What would be the correct interpretation of negative value in this case? In the example you have mentioned since both the vectors were same the value we received was -1.0. When i google the meaning of it certain blogs mentioned that it means that vector are similar but in opposite directions . This does not seem a correct interpretation as both vectors are same

Start Machine Learning

Never miss a tutorial:

Jason Brownlee

November 8, 2018 at 2:11 pm

#

in

twitter


f

envelope

rss

ty, I don't have material on this measure, perhaps this will help:
https://en.wikipedia.org/wiki/Cosine_similarity

Picked for you:



Your First Deep Learning Project in Python

w/ Jenna Ma


December 24, 2018 at 9:47 pm

#

This post helps me again. You are the best. 😊

I have a question when I write the custom metrics for my project. I want to define it as the MSE of predicted g and observed g. When I write a custom metric to calculate the MSE, I don't know how to make y_true represents the observed g. Did I make it clear?

Regression Tutorial with the Keras Deep Learning Library in Python




Multi-Class Classification Tutorial with the Keras

Jason Brownlee

December 25, 2018 at 7:21 am

#

It might be easier to write a custom function and evaluate model performance manually.



How to Save and Load Your Keras Deep Learning Model

Jenna Ma

December 26, 2018 at 5:40 pm

#

Loving The Tutorials? advice. I use the method you introduced in another post: <https://machinelearningmastery.com/implement-machine-learning-algorithm-performance-metrics-scratch-python/> The Deep Learning with Python EBook is where you'll find the Really Good stuff.

>> SEE WHAT'S INSIDE

Jason Brownlee

December 27, 2018 at 5:40 am

#

Well done!

Jimmy Tsao

December 25, 2018 at 8:25 pm

#

I'm trying to build my own accuracy function that checks if the output sequence is same as the true answer . For example if the true answer is " 0.2 0.4 0.6 0.8 ", either " 0.4 0.6 0.8 0.2 " or "0.8 0.6 0.4 0.2 " will be define as correct. Do you have any thoughts or recommendations?

Jason Brownlee

December 26, 2018 at 6:42 am

#

You would not use accuracy, you would use an error, such as MSE, MAE or RMSE.

Jimmy Tsao

December 26, 2018 at 2:27 pm

#

But I would like the out come be 1's and 0's not in the middle. Will that be possible?

Also merry Christmas, forgot that yesterday.

Jason Brownlee

December 27, 2018 at 5:39 am

#

You can round a floating point value to either 0/1.

Jenna Ma

December 27, 2018 at 7:20 pm

#

I did not find any post in your blog specifically focus on loss function, so I submit my question under this post. There are two output features in my model. I'm not sure how the loss function works. Whether the loss function returns the sum of two calculated errors or weighted sum or some other values? MSE = MSEa + MSEb ? MSE = 0.5*MSEa + 0.5*MSEb ? If it returns the weighted sum, can I define the weight? Thank you in advance.

Jason Brownlee

December 28, 2018 at 5:55 am

#

You can choose how to manage how to calculate loss on multiple outputs.

Adding a constant 1 or 0.5 does not make any difference in practice, I would imagine.

Ayhan

December 30, 2018 at 11:53 pm

#

Hi Dr.Brownlee, I want to define custom monitor metrics such as AUC for Early Stopping and ModelCheckpoint and other callbacks for monitor options and metrics for model.compile, What do i do ? I'm looking forward to your answer. Thanks,

Jason Brownlee

December 31, 2018 at 6:11 am

#

Create a list of callbacks and pass it to the "callbacks" argument on the fit() function.

Start Machine Learning

^

Start Machine Learning

You can master applied Machine Learning without math or fancy degrees. Find out how in this free and practical course.

Email Address

START MY EMAIL COURSE

<https://machinelearningmastery.com/custom-metrics-deep-learning-keras-python/>

11/21

Never miss a tutorial:

Solo January 7, 2019 at 3:23 pm #



Precision and Recall metrics have been removed from the latest version of keras, they cited that the metric was misleading, do you have any idea how to create a custom precision and recall metrics?

Picked for you:



Your First Deep Learning Project in Python
with Keras
Jason Brownlee January 8, 2019 at 6:44 am #

Yes, you can make predictions with your model then calculate the metrics with sklearn:



<https://scikit-learn.org/stable/modules/classes.html#module-sklearn.metrics>
Deep Learning Models in Python With Keras



Asal January 28, 2019 at 3:31 pm #
Regression Tutorial with the Keras Deep Learning Library in Python
Thank you so much for your tutorials.

I want to write a custom metric function. This function is PSNR (Peak signal-to-noise ratio) which is most commonly used to measure the quality of reconstruction of loss. PSNR is calculated based on the MSE results. So I was wondering if there is a way to write this PSNR function using the loss that is calculated in the fitting process. Or use that information to calculate PSNR.

Also, In your code:



[How to Save and Load Your Keras Deep Learning Model](#)
rm backend.sqrt(backend.mean(backend.square(y_pred - y_true), axis=-1))

What are the inputs of rmse during the training? How it is assigning y_true and y_pred?

Loving the Tutorials?

The [Deep Learning with Python](#) EBook is where you can find the [Really Good stuff](#)
The Really Good stuff January 29, 2019 at 6:08 am #

>> SEE WHAT'S INSIDE tion are the true y values and the predicted y values.

Asal January 30, 2019 at 9:22 am #

Thank you so much for your response, Jason.

This custom metric should return a tensor, right?

I had to use log10 in my computations. But then Keras only has log of e. (tf.keras.backend.log(x))

So I used math.log10 and I was getting an error in model.compile(). Here is the previous code:

```
def PSNR(y_true, y_pred):
    max_l = 1.0
    return 20*math.log10(max_l) - 10*math.log10( backend.mean( backend.square(y_pred - y_true),axis=-1))
```

Then, I thought I can use numpy to calculate the last line and then make a tensor of the result.

So I changed my code to:

```
def PSNR(y_true, y_pred):
    max_l = 1.0
    val = 20*math.log10(max_l) - 10*math.log10(np.mean( np.square(y_pred - y_true),axis=-1))
    newTensor = K.variable(value = val)
    return newTensor
```

I'm not getting any error. But is this the right way to do this?

Jason Brownlee January 30, 2019 at 2:40 pm #

Yes, you must operate upon tensors.

Sorry, I don't have the capacity to review/debug your approach.

Asal January 31, 2019 at 6:41 am #

Thank you, Jason.


Niraj March 1, 2019 at 2:04 am #

```
1 from numpy import array
2 from keras.models import Sequential
3 from keras.layers import Dense
4 from matplotlib import pyplot
5 # prepare sequence
6 X = array([0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0])
7 y = array([0, 0, 0, 0, 0, 1, 1, 1, 1, 1])
8 # create model
9 model = Sequential()
10 model.add(Dense(2, input_dim=1))
11 model.add(Dense(1, activation='sigmoid'))
12 model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['sparse_categorical_accuracy'])
13 # train model
14 history = model.fit(X, y, epochs=40, batch_size=len(X), verbose=2)
15 # plot metrics
16 acc = history.history['sparse_categorical_accuracy']
17 val_acc = history.history['val_acc']
18 epochs = range(1, len(acc) + 1)
19 pyplot.plot(epochs, acc, 'bo', label='Training acc')
20 pyplot.title('Training and validation accuracy')
21 pyplot.legend()
22 pyplot.tight_layout()
23 pyplot.show()
24
25 Epoch 39/40
26 - 0s - loss: 0.6356 - sparse_categorical_accuracy: 0.5000
27 Epoch 40/40
28 - 0s - loss: 0.6353 - sparse_categorical_accuracy: 0.5000
29
30 KeyError                                Traceback (most recent call last)
31 in C
32     15 # plot metrics
33     16 acc = history.history['sparse_categorical_accuracy']
```

Start Machine Learning


```
Ne 34 --> 17 val_acc = history.history['val_acc']
35 miss a tutorial: range(1, len(acc) + 1)
36 19 pyplot.plot(epochs,acc, 'bo', label='Training acc')
37
38 KeyError: 'val_acc'
```

How do I resolve this error message? Please help!
Picked for you:




Jason Brownlee March 1, 2019 at 6:25 am #

Your First Deep Learning Project in Python with Keras Step-By-Step




Krishna March 4, 2019 at 11:17 pm #

How to Grid Search Hyperparameters for Deep Learning Models in Python With Keras




Jason Brownlee March 5, 2019 at 6:40 am #

Regression Tutorial with the Keras Deep Learning Library



Jason Brownlee March 5, 2019 at 6:40 am #

Multi-Class Classification Tutorial with the Keras Deep Learning Library



Jason Brownlee March 5, 2019 at 6:40 am #

How to Use Custom Metrics in Keras

In order to access val_acc you must fit the model with a validation dataset. e.g. set validation_data=(...) in the call to model.fit(...)

can a large value for mse can be given? the system be tested for convergence

This is a common question that I answer here: <https://machinelearningmastery.com/faq/single-faq/how-to-know-if-a-model-has-good-performance>

Start Machine Learning

You can master applied Machine Learning without math or fancy degrees. Find out how in this free and practical course.

Email Address

START MY EMAIL COURSE

Lowering the Tutorials? #

The *Deep Learning with Python* EBook is where you'll find the *Really Good* stuff.

Is there ever a limit to number of epochs? In my data set (regression) the more epochs the better the model keeps performing... Even past 500... Is anything over 2000 epochs odd??

>> SEE WHAT'S INSIDE

Jason Brownlee March 6, 2019 at 2:44 pm #

Not really.

ben March 7, 2019 at 1:13 am #

In regression... Ideally when should one stop adding epochs? Is it possible to verify just thru an RSME plot?

Jason Brownlee March 7, 2019 at 6:54 am #

When the model no longer improves on the holdout validation dataset.

R April 22, 2019 at 12:22 am #

Hi Jason,

In the codes of Custom Metrics in Keras part, you defined the rmse function as follow:

```
def rmse(y_true, y_pred):
    return backend.sqrt(backend.mean(backend.square(y_pred - y_true), axis=-1))
```

Why is it necessary to write axis=-1? I don't understand what axis=-1 means here. I deleted axis=-1 from the function in my codes but it is still OK to run?

Jason Brownlee April 22, 2019 at 6:25 am #

It is explicitly specifying to calculate the error across the last dimension, normally this is samples, but for encoder-decoder lstrms this will be time steps.

Belgian_student May 15, 2019 at 10:17 am #

Hello Mr. Brownlee,

-> Thanks for this tutorial, you helped me a lot.

For my thesis, I did a regression cnn in keras using the four metrics you present here to be interesting for regression.

I've this question with respect to the cosine_proximity-metric: if the vectors in nominator and denominator of the underlying formula are 1-dimensional vectors (=being just the real valued true and predicted label), then the metric will always resolve to 1?

- Why is its use then so interesting for regression networks, or maybe networks with multiple regressed output are intended here?

- Have you got any idea how its value could be '-1' instead of '+1' when both the true and predicted label are positive?

Kind Regards from Belgium

Jason Brownlee May 15, 2019 at 2:45 pm #

Sorry, I can't give you good off the cuff about the cosine similarity metric.

I hope to cover it in the future.

Elena June 16, 2019 at 12:05 am #


Hi!

I am trying to train a recurrent neural network implemented using Keras and mean square error as loss function. Is it possible to have a loss greater than 1 and the model generated by the network to work as expected?

Never miss a tutorial:
Jason Brownlee June 16, 2019 at 7:13 am #

REPLY


Picked for you:



shiva July 26, 2019 at 12:05 am #

Yoshua Bengio's Deep Learning in Python with Keras Step-By-Step
When i try to use a model saved using rmse as metric.
During loading the model
I get the following error:
ValueError: Unknown metric function:rmse

What is the best metric for timeseries data?
My model with MSE is either good in capturing higher signals or either fails to capture low signals..
I want a better metric which would preserve correlation and MSE together..
thank you



Multi-Class Classification Tutorial with the Keras Deep Learning Library

Jason Brownlee July 25, 2019 at 2:11 pm #

How to Save and Load Your Keras Deep Learning Model
Good question, you must provide a dict to the load_model() function that indicates what the rmse function means.
For example, and assuming the rmse function is defined:

```
1 model = load_model('model.h4', custom_objects={'rmse':rmse})
```

Loving the Tutorials?

The Deep Learning with Python EBook is where you'll find the Really Good stuff.
shiva July 26, 2019 at 12:39 am #

>> SEE WHAT'S INSIDE

I copy but i still have an error.
IndexError: tuple index out of range
C:\ProgramData\Anaconda3\lib\site-packages\numpy\core_methods.py in _count_reduce_items(arr, axis)
53 items = 1
54 for ax in axis:
--> 55 items *= arr.shape[ax]
56 return items
57
IndexError: tuple index out of range

REPLY

Jason Brownlee July 26, 2019 at 8:26 am #

Sorry to hear that. I don't have any good ideas.
Perhaps try searching/posting stackoverflow?

REPLY

shiva July 26, 2019 at 12:05 am #

What is the best metric for timeseries data?
My model with MSE is either good in capturing higher signals or either fails to capture low signals..
I want a better metric which would preserve correlation and MSE together..
REPLY

Jason Brownlee July 26, 2019 at 8:26 am #

I find RMSE or MAPE useful.
REPLY

Anand August 20, 2019 at 1:05 am #

Very informative blog. But can you please tell me how to use recall as a metric.
REPLY

Jason Brownlee August 20, 2019 at 6:27 am #

Just plug it in to the above examples.
REPLY

Alex September 7, 2019 at 10:15 am #

Hello Sir.
Please if I've normalized my dataset (X and Y), with MinMaxScaler for example, and if I'm using MSE or RMSE for loss and/or for metrics, the results expected (mse and rmse) are also normalized, right?
How can I get the "real" MSE and RMSE of the original data (X and Y) denormalized?
Thanks much in advance.
REPLY

Jason Brownlee September 8, 2019 at 5:12 am #

Correct.
You can get real error by inverting the transform on the predictions first, then calculating error metrics. The objects typically offer an inverse_transform() function.
REPLY

Alex September 12, 2019 at 6:36 am #

Ok. Right. I did it. Thanks! One more question please...
But how about if I, let's say, normalize X and standardize Y, or vice-versa.
REPLY

Start Machine Learning

Start Machine Learning

You can master applied Machine Learning without math or fancy degrees. Find out how in this free and practical course.
Email Address
START MY EMAIL COURSE

https://machinelearningmastery.com/custom-metrics-deep-learning-keras-python/

14/21

Never miss a tutorial:

When inverting the transformation on the predictions [predict(X_test) = Y_pred], which scaler should I use to get the "real" Y_pred inversely transformed?

The inverse of normalized or the inverse of standardized?



Thanks much in advance.

Picked for you:



Your First Deep Learning Project in Python
with Keras Step-By-Step

Jason Brownlee September 12, 2019 at 1:46 pm #

REPLY ↩



You invert the transforms applied to y, in the reverse order in which they were applied.
This may give you some ideas:

Keras <https://machinelearningmastery.com/machine-learning-data-transforms-for-time-series-forecasting/>



Regression Tutorial with the Keras Deep
Learning Library in Python

benjamin appiah October 12, 2019 at 5:09 pm #



I tried using a custom loss function but always fall into errors.

Multi-Class Classification Tutorial with the

Keras Deep Learning Library in Python

How to Save and Load Your Keras Deep

Learning Model

How to Save and Load Your Keras Deep

Learning Model

How to Save and Load Your Keras Deep

Learning Model

How to Save and Load Your Keras Deep

Learning Model

How to Save and Load Your Keras Deep

Learning Model

How to Save and Load Your Keras Deep

Learning Model

How to Save and Load Your Keras Deep

Learning Model

How to Save and Load Your Keras Deep

Learning Model

How to Save and Load Your Keras Deep

Learning Model

How to Save and Load Your Keras Deep

Learning Model

How to Save and Load Your Keras Deep

Learning Model

How to Save and Load Your Keras Deep

Learning Model

How to Save and Load Your Keras Deep

Learning Model

How to Save and Load Your Keras Deep

Learning Model

How to Save and Load Your Keras Deep

Learning Model

How to Save and Load Your Keras Deep

Learning Model

How to Save and Load Your Keras Deep

Learning Model

How to Save and Load Your Keras Deep

Learning Model

How to Save and Load Your Keras Deep

Learning Model

How to Save and Load Your Keras Deep

Learning Model

How to Save and Load Your Keras Deep

Learning Model

How to Save and Load Your Keras Deep

Learning Model

How to Save and Load Your Keras Deep

Learning Model

How to Save and Load Your Keras Deep

Learning Model

How to Save and Load Your Keras Deep

Learning Model

How to Save and Load Your Keras Deep

Learning Model

How to Save and Load Your Keras Deep

Learning Model

How to Save and Load Your Keras Deep

Learning Model

How to Save and Load Your Keras Deep

Learning Model

How to Save and Load Your Keras Deep

Learning Model

How to Save and Load Your Keras Deep

Learning Model

How to Save and Load Your Keras Deep

Learning Model

How to Save and Load Your Keras Deep

Learning Model

How to Save and Load Your Keras Deep

Learning Model

How to Save and Load Your Keras Deep

Learning Model

How to Save and Load Your Keras Deep

Learning Model

How to Save and Load Your Keras Deep

Learning Model

How to Save and Load Your Keras Deep

Learning Model

How to Save and Load Your Keras Deep

Learning Model

How to Save and Load Your Keras Deep

Learning Model

How to Save and Load Your Keras Deep

Learning Model

How to Save and Load Your Keras Deep

Learning Model

How to Save and Load Your Keras Deep

Learning Model

How to Save and Load Your Keras Deep

Learning Model

How to Save and Load Your Keras Deep

Learning Model

How to Save and Load Your Keras Deep

Learning Model

How to Save and Load Your Keras Deep

Learning Model

How to Save and Load Your Keras Deep

Learning Model

```
d(x,y) = square [Transpose(x-y) * Inverse(S) * (x-y)]
mahalanobis distance for generalized squared interpoint distance" for its squared value(3)] can also be defined as a dissimilarity measure between two random vectors
matrix S.
d(x,y) = square [Transpose(x-y) * Inverse(S) * (x-y)]
s://en.wikipedia.org/wiki/Mahalanobis_distance)
Learning Model
n_classes = 4
n_samples=800
X, y = make_classification(n_samples=n_samples, n_features=20, n_informative=4, n_redundant=0, n_classes=n_classes, n_clusters_per_class=2)
y = to_categorical(y, n_classes)
Xtrainb, testXb, ytrainb, ytestb = train_test_split(X, y, test_size = 0.3, random_state=42)
x_trainb = ytrainb.reshape((ytrainb.shape[0], Xtrainb.shape[1], 1))
Xtestb = ytestb.reshape((ytestb.shape[0], Xtestb.shape[1], 1))
densesize = 4
input_datab = Input(shape=(Xtrainb.shape[1],1))
epochs = 10
batch_size = 32

#####
def mahalanobis(y_true, y_pred):
    x_minus_mn_with_transpose = K.transpose(y_true - y_pred)
    Covariance = covr1(y_true, y_pred)
    inv_covmat = tf.linalg.inv(Covariance)
    x_minus_mn = y_true - y_pred
    left_term = K.dot(x_minus_mn, inv_covmat)
    D_square = K.dot(left_term, x_minus_mn_with_transpose)
    return D_square

def covr1(y_true, y_pred):
    #x_mean = K.mean(y_true)
    #y_mean = K.mean(y_pred)
    Cov_numerator = K.sum(((y_true - y_pred)*(y_true - y_pred)))
    Cov_denominator = len(Xtrainb)-1
    Covariance = (Cov_numerator / Cov_denominator)
    return Covariance

conv1= Conv1D(filters=80, kernel_size=2, padding='same', input_dim=Xtrainb.shape[1])(input_datab)
maxpool = MaxPooling1D(pool_size=3, stride=3 )(conv1)
conv2= Conv1D(filters=50, kernel_size=2, padding='same', input_dim=Xtrainb.shape[1])(maxpool)
maxpool = MaxPooling1D(pool_size=3, stride=3)(conv2)
flatten = Flatten()(maxpool)
dense = Dense(84, activation='relu')(flatten)
dense = Dense(1024, activation='relu')(dense)
dense = Dense(densesize, activation='softmax')(dense)
model = Model(inputs=[input_datab], outputs=[dense])
model.compile(loss= mahalanobis, optimizer='adam', metrics=['acc'])
hist = model.fit(x_trainb, ytrainb, validation_data=(Xtestb, ytestb), epochs=epochs, batch_size=batch_size)
```

Loving the Tutorials?

The Deep Learning with Python EBook is

available for free on the

Deep Learning with Python EBook is

available for free on the

Deep Learning with Python EBook is

available for free on the

Deep Learning with Python EBook is

available for free on the

Deep Learning with Python EBook is

available for free on the

Deep Learning with Python EBook is

available for free on the

Deep Learning with Python EBook is

available for free on the

Deep Learning with Python EBook is

available for free on the

Deep Learning with Python EBook is

available for free on the

Deep Learning with Python EBook is

available for free on the

Deep Learning with Python EBook is

available for free on the

Deep Learning with Python EBook is

available for free on the

Deep Learning with Python EBook is

available for free on the

Deep Learning with Python EBook is

available for free on the

Deep Learning with Python EBook is

available for free on the

Deep Learning with Python EBook is

available for free on the

Deep Learning with Python EBook is

available for free on the

Deep Learning with Python EBook is

available for free on the

Deep Learning with Python EBook is

available for free on the

Deep Learning with Python EBook is

available for free on the

Deep Learning with Python EBook is

available for free on the

Deep Learning with Python EBook is

available for free on the

Deep Learning with Python EBook is

available for free on the

Deep Learning with Python EBook is

available for free on the

Deep Learning with Python EBook is

available for free on the

Deep Learning with Python EBook is

available for free on the

Deep Learning with Python EBook is

available for free on the

Deep Learning with Python EBook is

available for free on the

Deep Learning with Python EBook is

available for free on the

Deep Learning with Python EBook is

available for free on the

Deep Learning with Python EBook is

available for free on the

Deep Learning with Python EBook is

available for free on the

Deep Learning with Python EBook is

available for free on the

Deep Learning with Python EBook is

available for free on the

Deep Learning with Python EBook is

available for free on the

Deep Learning with Python EBook is

available for free on the

Deep Learning with Python EBook is

available for free on the

Deep Learning with Python EBook is

available for free on the

Deep Learning with Python EBook is

available for free on the

Deep Learning with Python EBook is

available for free on the

Deep Learning with Python EBook is

available for free on the

Deep Learning with Python EBook is

available for free on the

Deep Learning with Python EBook is

available for free on the

Deep Learning with Python EBook is

available for free on the

Jason Brownlee October 13, 2019 at 8:28 am #

REPLY ↩

Sorry to hear that, I have some suggestions here that might help:

<https://machinelearningmastery.com/faq/single-faq/can-you-read-review-or-debug-my-code>

Dawit December 14, 2019 at 11:09 pm #

REPLY ↩

Hello Sir,

I was developing MLPRegressor model like...

#create a model

nn=MLPRegressor(hidden_layer_sizes=(2, 1.),activation='logistic',max_iter=2000,solver='adam',learning_rate_init=0.1,momentum=0.7,early_stopping=True,

validation_fraction=0.15,)

history = nn.fit(X_train, y_train,)

how can I plot mape, r^2 and how can I predict for new samples. I was scaled my data using minmax scaler???

Jason Brownlee December 15, 2019 at 6:06 am #

REPLY ↩

Make a prediction on the dataset then plot the real y values vs the predicted y values.

If you are using scikit-learn, not keras, then this will help you make a prediction:

<https://machinelearningmastery.com/make-predictions-scikit-learn/>

Frank Tang February 28, 2020 at 10:24 am #

REPLY ↩


Dear Prof. Brownlee:

I try the following code:

Start Machine Learning


Never miss a tutorial. Squared Error are [6.21791493e-02 3.92977809e-02 2.16430749e-02 9.21505186e-03 2.01369724e-03 3.90194594e-05 3.29101280e-03 1.17696773e-02 2.1750775e-02 1.4070447e-02] Mean Squared Error are 0.021933054033792435 Root Mean Squared Error is 0.14809812299213124

Picked for you: Notice the evaluate return 0.1278020143508911 instead of the correct 0.14809812299213124




Your First Deep Learning Project in Python with Keras Step-By-Step

Jason Brownlee February 29, 2020 at 7:08 am #




How to Grid Search Hyperparameters for Deep Learning Models in Python With Keras

Thanks, I will investigate.




Regression Tutorial with the Keras Deep Learning Library in Python

How can I get different components of the loss function if I am using model.train_on_batch instead of model.fit? I have seen that model.train_on_batch returns components of loss functions?



Multi-Class Classification Tutorial with the Keras Deep Learning Library



How to Save and Load Your Keras Deep Learning Model

Jason Brownlee March 19, 2020 at 6:30 am #

I'm not sure off hand, perhaps one of these resources will help:

<https://machinelearningmastery.com/get-help-with-keras/>

Loving the Tutorials?

Start Machine Learning

You can master applied Machine Learning without math or fancy degrees. Find out how in this free and practical course.

START MY EMAIL COURSE

The Deep Learning with Python EBook is where you find the Really Good stuff.

>> SEE WHAT'S INSIDE

```
# build encoder model
def encoder_model(inputs):
    x1 = Dense(intermediate_dim_1, activation='relu')(inputs)
    x2 = Dense(intermediate_dim_2, activation='relu')(x1)
    x3 = Dense(intermediate_dim_3, activation='relu')(x2)
    x4 = Dense(intermediate_dim_4, activation='relu')(x3)
    z_mean_encoded = Dense(latent_dim, name='z_mean')(x4)
    z_log_var_encoded = Dense(latent_dim, name='z_log_var')(x4)

    # instantiate encoder model
    encoder = Model(inputs, [z_mean_encoded, z_log_var_encoded], name='encoder')
    return encoder, z_mean_encoded, z_log_var_encoded

# build decoder model
def decoder_model():
    latent_inputs = Input(shape=(latent_dim,), name='z_sampling')
    x4 = Dense(intermediate_dim_4, activation='relu')(latent_inputs)
    x3 = Dense(intermediate_dim_3, activation='relu')(x4)
    x2 = Dense(intermediate_dim_2, activation='relu')(x3)
    x1 = Dense(intermediate_dim_1, activation='relu')(x2)
    outputs = Dense(original_dim)(x1)

    # instantiate decoder model
    decoder = Model(latent_inputs, outputs, name='decoder')
    return decoder

def recon_loss(inputs, outputs):
    reconstruction_loss = mse(inputs, outputs)
    return K.mean(reconstruction_loss)

def kl_loss():
    kl_loss = 1 + z_log_var_encoded - K.square(z_mean_encoded) - K.exp(z_log_var_encoded)
    kl_loss = K.sum(kl_loss, axis=-1)
    kl_loss *= -0.5
    return K.mean(kl_loss)

# # reconstruction_loss *=
# kl_loss = 1 + z_log_var_encoded - K.square(z_mean_encoded) - K.exp(z_log_var_encoded)
# kl_loss = K.sum(kl_loss, axis=-1)
# kl_loss *= -0.5
# kl_loss_metric = kl_loss
# kl_loss *= beta
# vae_loss = K.mean(reconstruction_loss + kl_loss)

def total_loss(inputs, outputs, z_mean_encoded, z_log_var_encoded, beta):
    reconstruction_loss = mse(inputs, outputs)
    kl_loss = 1 + z_log_var_encoded - K.square(z_mean_encoded) - K.exp(z_log_var_encoded)
    kl_loss = K.sum(kl_loss, axis=-1)
    kl_loss *= -0.5
    kl_loss *= beta
    return K.mean(reconstruction_loss + kl_loss)

def sampling(args):
    """Reparameterization trick by sampling fr an isotropic unit Gaussian.
    # Arguments:
    args (tensor): mean and log of variance of Q(z|X)
    # Returns:
    z (tensor): sampled latent vector
    """
    z_mean, z_log_var = args
    batch = K.shape(z_mean)[0]
    dim = K.int_shape(z_mean)[1] # Returns the shape of tensor or variable as a tuple of int or None entries.
    # by default, random_normal has mean=0 and std=1.0
    epsilon = K.random_normal(shape=(batch, dim))
    return z_mean + K.exp(0.5 * z_log_var) * epsilon

if __name__ == '__main__':
    x_trn, x_val, y_trn, y_val = train_test_split(Cp_inputs, X_all, test_size=0.2, shuffle=True, random_state=0)
    original_dim = x_trn.shape[1]
    x_trn = np.reshape(x_trn, [-1, original_dim])
    x_val = np.reshape(x_val, [-1, original_dim])
```

Start Machine Learning

Never miss a tutorial.

input_shape = (original_dim,)

inputs = Input(shape=input_shape, name='encoder_input')

Define Intermediate Layer Dimension and Latent layer Dimension

intermediate_dim_1 = 128

intermediate_dim_2 = 256

intermediate_dim_3 = 128

intermediate_dim_4 = 64

Create the encoder model

encoder = Sequential([

Dense(intermediate_dim_1, activation='relu'),

Dense(intermediate_dim_2, activation='relu'),

Dense(intermediate_dim_3, activation='relu'),

Dense(intermediate_dim_4, activation='relu'),

])

Create the decoder model

decoder = Sequential([

Dense(intermediate_dim_4, activation='relu'),

Dense(intermediate_dim_3, activation='relu'),

Dense(intermediate_dim_2, activation='relu'),

Dense(intermediate_dim_1, activation='relu'),

])

Create the VAE model

vae = Model(inputs, outputs, name='vae')

Compile the VAE model

vae.compile(optimizer=Adam(), metrics=[recon_loss, latent_loss])

Train the VAE model

vae.fit(x_train, y_train, epochs=epochs, batch_size=batch_size, validation_data=(x_val, None), verbose = 2)

input_shape = (original_dim,)

inputs = Input(shape=input_shape, name='encoder_input')

Define Intermediate Layer Dimension and Latent layer Dimension

intermediate_dim_1 = 128

intermediate_dim_2 = 256

intermediate_dim_3 = 128

intermediate_dim_4 = 64

Create the encoder model

encoder = Sequential([

Dense(intermediate_dim_1, activation='relu'),

Dense(intermediate_dim_2, activation='relu'),

Dense(intermediate_dim_3, activation='relu'),

Dense(intermediate_dim_4, activation='relu'),

])

Create the decoder model

decoder = Sequential([

Dense(intermediate_dim_4, activation='relu'),

Dense(intermediate_dim_3, activation='relu'),

Dense(intermediate_dim_2, activation='relu'),

Dense(intermediate_dim_1, activation='relu'),

])

Create the VAE model

vae = Model(inputs, outputs, name='vae')

Compile the VAE model

vae.compile(optimizer=Adam(), metrics=[recon_loss, latent_loss])

Train the VAE model

vae.fit(x_train, y_train, epochs=epochs, batch_size=batch_size, validation_data=(x_val, None), verbose = 2)

Start Machine Learning

You can master applied Machine Learning without math or fancy degrees. Find out how in this free and practical course.

Email Address

START MY EMAIL COURSE

Result : **Loving the Tutorials?**

Epoch 1/10

– 0s – loss: 3.42770 – val_loss: 4.9581

Epoch 2/10

– 0s – loss: 4.1527 – val_loss: 2.4654

Epoch 3/10

– 0s – loss: 3.2343 – val_loss: 2.7032

Epoch 4/10

– 0s – loss: 2.5479 – val_loss: 2.5234

Epoch 5/10

– 0s – loss: 2.3551 – val_loss: 2.2926

Epoch 6/10

– 0s – loss: 2.2032 – val_loss: 2.1937

Epoch 7/10

– 0s – loss: 1.9983 – val_loss: 2.0159

Epoch 8/10

– 0s – loss: 1.8385 – val_loss: 1.6428

Epoch 9/10

– 0s – loss: 1.6508 – val_loss: 1.5881

Epoch 10/10

– 0s – loss: 1.5189 – val_loss: 1.4624

I am trying to make VAE model, but it does not give any metric values, which were defined as [recon_loss, latent_loss]

Can you solve this issue?

Jason Brownlee April 2, 2020 at 5:43 am #

REPLY

This is a common question that I answer here:
<https://machinelearningmastery.com/faq/single-faq/can-you-read-review-or-debug-my-code>

kiki April 7, 2020 at 6:15 pm #

REPLY

Hi Jason, can I find the accuracy of keras regression problem? From your notes, for keras regression problem only mse,rmse,mae. Is it possible for me to find the accuracy of this method?

Jason Brownlee April 8, 2020 at 7:48 am #

REPLY

No, see this:
<https://machinelearningmastery.com/faq/single-faq/how-do-i-calculate-accuracy-for-regression>

kiki April 8, 2020 at 12:30 pm #

REPLY

Thanks Jason 😊

Jason Brownlee April 8, 2020 at 1:20 pm #

REPLY

You're welcome!

kiki April 9, 2020 at 6:12 pm #

REPLY

Hi Jason, I want to ask you how to know whether the model provide a good performance for regression? Because previously u said that we cannot know the accuracy of regression. Is it by their loss mse,mae and rmse to decide the model has the good performance? I mean if the loss of mse is below than 1, then the model are good?

Jason Brownlee April 10, 2020 at 8:23 am #

REPLY

Good question, see this:
<https://machinelearningmastery.com/faq/single-faq/how-to-know-if-a-model-has-good-performance>

Start Machine Learning

<https://machinelearningmastery.com/custom-metrics-deep-learning-keras-python/>

18/21


Never miss a tutorial


Hi,


I have Sub-Classed the Metric class to create a custom precision metric. Everything looks fine; I mean there is no run-time error. But I suspect there is something wrong when I see the precision scores logging in the output


Picked for you:


Here is my code and its output:

[Your First Deep Learning Project in Python](#)

[How to Grid Search Hyperparameters for Precision Metrics \(Keras metrics.Metric\)](#)

[Regression Tutorial with the Keras Deep Learning Library in Python](#)

[Multi-Class Classification Tutorial with the Keras Deep Learning Library](#)

[How to Save and Load Your Keras Deep Learning Model](#)

```
all_train_data, X_train_10, y_train_10, X_test_10, y_test_10) = keras.datasets.cifar10.load_data()

X_train_10 = X_train_10 / 255.
X_test_10 = X_test_10 / 255.

class PrecisionMetric(keras.metrics.Metric):
    def __init__(self, name = 'precision', **kwargs):
        super(PrecisionMetric, self).__init__(**kwargs)
        self.add_weight('tp', initializer = 'zeros')
        self.add_weight('fp', initializer = 'zeros')

    def update_state(self, y_true, y_pred):
        y_true = tf.cast(y_true, tf.bool)
        y_pred = tf.cast(y_pred, tf.bool)
        p = tf.logical_and(tf.equal(y_true, True), tf.equal(y_pred, True))
        false_p = tf.logical_and(tf.equal(y_true, False), tf.equal(y_pred, True))
        self.tp.assign_add(tf.reduce_sum(tf.cast(true_p, self.dtype)))
        self.fp.assign_add(tf.reduce_sum(tf.cast(false_p, self.dtype)))
        self.reset_states(self)

    def result(self):
        self.tp.assign(0)
        self.fp.assign(0)

    def result(self):
        return self.tp / (self.tp + self.fp)

# The Deep Learning with Python EBook is
# where you will find the Really Good stuff.
model = keras.models.Sequential([
    >> SEE WHAT'S INSIDE

model.add(keras.layers.Flatten(input_shape = np.array(X_train_10.shape[1: ])))
for _ in range(2):
    model.add(keras.layers.Dense(50, activation = 'elu', kernel_initializer = 'he_normal'))
    model.add(keras.layers.Dense(1, activation = 'sigmoid'))

loss = keras.losses.binary_crossentropy
optimizer = keras.optimizers.SGD()

model.compile(loss = loss, optimizer = optimizer, metri

# To make it binary classification
y_train_5 = (y_train_10 == 5)
y_test_5 = (y_test_10 == 5)

history = model.fit(X_train_10, y_train_5, epochs = 5)

Epoch 1/5
1563/1563 [=====] - 5s 3ms/step - loss: 0.2954
Epoch 2/5
1563/1563 [=====] - 4s 3ms/step - loss: 0.2779
Epoch 3/5
1563/1563 [=====] - 4s 2ms/step - loss: 0.2701
Epoch 4/5
1563/1563 [=====] - 3s 2ms/step - loss: 0.2660
Epoch 5/5
1563/1563 [=====] - 3s 2ms/step - loss: 0.2629
```

REPLY

Start Machine Learning

You can master applied Machine Learning without math or fancy degrees. Find out how in this free and practical course.

START MY EMAIL COURSE

Jason Brownlee

December 23, 2020 at 5:36 am

#

Thanks for sharing.

Sorry, I don't have the capacity to review/debug your code.

REPLY

JG

January 17, 2021 at 11:00 pm

#

Hi Jason,

interesting tutorial !

when using proper (custom) metrics (e.g. 'rmse') after saving the keras model (via .save method()) when you want to load again the model (via load_model() method), it give you an error because it does not understand your own defined 'rmse' metric... how can we solve the keras loading?

thks

REPLY

Jason Brownlee

January 18, 2021 at 6:08 am

#

Thanks!

Good question, you will need to have the function defined when loading the model and specify the function via the custom_objects argument to the load_model() function.

https://keras.io/api/models/model_saving_apis/

... custom_objects: Optional dictionary mapping names (strings) to custom classes or functions to be considered during deserialization.

REPLY

JG

January 18, 2021 at 11:40 pm

#

I see:

model = load_model('model.h5', custom_objects={'rmse':rmse})

thks you very much Jason!

REPLY

Jason Brownlee

January 19, 2021 at 6:38 am

#

Well done! You're welcome.

Start Machine Learning

Never miss a tutorial:
Resam February 20, 2021 at 11:43 pm #


in

hi,

f

Im trying to use mean absolute precentage error and i use loss: 'mse' and the mape results are around 600 and 800 what's the problem?

Picked for you:



Your First Deep Learning Project in Python with Keras Deep Learning Library

Jason Brownlee February 21, 2021 at 6:14 am #

Perhaps the model is a bad fit for your data?


Perhaps you need to use a different model configuration?

How to Grid Search Hyperparameters for Deep Learning Models in Python

Perhaps you need to use a different model?


Perhaps you need to use data preparation methods?

Perhaps your prediction problem is really hard?




Regression Tutorial with the Keras Deep Learning Library in Python

Leave a Reply



Multi-Class Classification Tutorial with the Keras Deep Learning Library



How to Save and Load Your Keras Deep Learning Model

Loving the Tutorials?


The Deep Learning with Python EBook is where you'll find the **Really Good** stuff.

>> SEE WHAT'S INSIDE

(will not be published) (required)

Website

SUBMIT COMMENT



Welcome!

I'm Jason Brownlee PhD and I help developers get results with machine learning.

[Read more](#)

Start Machine Learning

×

You can master applied Machine Learning **without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE

© 2021 Machine Learning Mastery Pty. Ltd. All Rights Reserved.

[Linkedin](#) | [Twitter](#) | [Facebook](#) | [Newsletter](#) | [RSS](#)

[Privacy](#) | [Disclaimer](#) | [Terms](#) | [Contact](#) | [Sitemap](#) | [Search](#)

Start Machine Learning

↑

https://machinelearningmastery.com/custom-metrics-deep-learning-keras-python/

21/21