

K

Keras

About Keras

Getting started

Developer guides

Keras API reference

Models API

Layers API

Callbacks API

Data preprocessing

Optimizers

Metrics

Losses

Built-in small datasets

Keras Applications

Utilities

Keras Tuner

Code examples

Why choose Keras?

Community & governance

Contributing to Keras

KerasTuner

Metrics

A metric is a function that is used to judge the performance of your model.

Metric functions are similar to loss functions, except that the results from evaluating a metric are not used when training the model. Note that you may use any loss function as a metric.

Available metrics

Accuracy metrics

- [Accuracy class](#)
- [BinaryAccuracy class](#)
- [CategoricalAccuracy class](#)
- [TopKCategoricalAccuracy class](#)
- [SparseTopKCategoricalAccuracy class](#)

Probabilistic metrics

- [BinaryCrossentropy class](#)
- [CategoricalCrossentropy class](#)
- [SparseCategoricalCrossentropy class](#)
- [KLDivergence class](#)
- [Poisson class](#)

Regression metrics

- [MeanSquaredError class](#)
- [RootMeanSquaredError class](#)
- [MeanAbsoluteError class](#)
- [MeanAbsolutePercentageError class](#)
- [MeanSquaredLogarithmicError class](#)
- [CosineSimilarity class](#)
- [LogCoshError class](#)

Classification metrics based on True/False positives & negatives

- [AUC class](#)
- [Precision class](#)
- [Recall class](#)
- [TruePositives class](#)
- [TrueNegatives class](#)
- [FalsePositives class](#)
- [FalseNegatives class](#)
- [PrecisionAtRecall class](#)
- [SensitivityAtSpecificity class](#)
- [SpecificityAtSensitivity class](#)

Image segmentation metrics

- [MeanIoU class](#)

Hinge metrics for "maximum-margin" classification

- [Hinge class](#)
- [SquaredHinge class](#)
- [CategoricalHinge class](#)

Usage with `compile()` & `fit()`

The `compile()` method takes a `metrics` argument, which is a list of metrics:

```
model.compile(
    optimizer='adam',
    loss='mean_squared_error',
    metrics=[
        metrics.MeanSquaredError(),
        metrics.AUC(),
    ]
)
```

Metric values are displayed during `fit()` and logged to the `History` object returned by `fit()`. They are also returned by `model.evaluate()`.

Note that the best way to monitor your metrics during training is via [TensorBoard](#).

To track metrics under a specific name, you can pass the `name` argument to the metric constructor:

```
model.compile(
    optimizer='adam',
    loss='mean_squared_error',
    metrics=[
        metrics.MeanSquaredError(name='my_mse'),
        metrics.AUC(name='my_auc'),
    ]
)
```

All built-in metrics may also be passed via their string identifier (in this case, default constructor argument values are used, including a default metric name):

```
model.compile(
    optimizer='adam',
    loss='mean_squared_error',
    metrics=[
        'MeanSquaredError',
        'AUC',
    ]
)
```

Standalone usage

Unlike losses, metrics are stateful. You update their state using the `update_state()` method, and you query the scalar metric result using the `result()` method:

```
m = tf.keras.metrics.AUC()
m.update_state([0, 1, 1, 1], [0, 1, 0, 0])
print('Intermediate result:', float(m.result()))

m.update_state([1, 1, 1, 1], [0, 1, 1, 0])
print('Final result:', float(m.result()))
```

The internal state can be cleared via `metric.reset_states()`.

Here's how you would use a metric as part of a simple custom training loop:

```

accuracy = tf.keras.metrics.CategoricalAccuracy()
loss_fn = tf.keras.losses.CategoricalCrossentropy(from_logits=True)
optimizer = tf.keras.optimizers.Adam()

# Iterate over the batches of a dataset.
for step, (x, y) in enumerate(dataset):
    with tf.GradientTape() as tape:
        logits = model(x)
        # Compute the loss value for this batch.
        loss_value = loss_fn(y, logits)

    # Update the state of the `accuracy` metric.
    accuracy.update_state(y, logits)

    # Update the weights of the model to minimize the loss value.
    gradients = tape.gradient(loss_value, model.trainable_weights)
    optimizer.apply_gradients(zip(gradients, model.trainable_weights))

    # Logging the current accuracy value so far.
    if step % 100 == 0:
        print('Step:', step)
        print('Total running accuracy so far: %.3f' % accuracy.result())

```

Creating custom metrics

As simple callables (stateless)

Much like loss functions, any callable with signature `metric_fn(y_true, y_pred)` that returns an array of losses (one of sample in the input batch) can be passed to `compile()` as a metric. Note that sample weighting is automatically supported for any such metric.

Here's a simple example:

```

def my_metric_fn(y_true, y_pred):
    squared_difference = tf.square(y_true - y_pred)
    return tf.reduce_mean(squared_difference, axis=-1) # Note the `axis=-1`

model.compile(optimizer='adam', loss='mean_squared_error', metrics=[my_metric_fn])

```

In this case, the scalar metric value you are tracking during training and evaluation is the average of the per-batch metric values for all batches seen during a given epoch (or during a given call to `model.evaluate()`).

As subclasses of `Metric` (stateful)

Not all metrics can be expressed via stateless callables, because metrics are evaluated for each batch during training and evaluation, but in some cases the average of the per-batch values is not what you are interested in.

Let's say that you want to compute AUC over a given evaluation dataset: the average of the per-batch AUC values isn't the same as the AUC over the entire dataset.

For such metrics, you're going to want to subclass the `Metric` class, which can maintain a state across batches. It's easy:

- Create the state variables in `__init__`
- Update the variables given `y_true` and `y_pred` in `update_state()`
- Return the metric result in `result()`
- Clear the state in `reset_states()`

Here's a simple example computing binary true positives:

```

class BinaryTruePositives(tf.keras.metrics.Metric):

    def __init__(self, name='binary_true_positives', **kwargs):
        super(BinaryTruePositives, self).__init__(name=name, **kwargs)
        self.true_positives = self.add_weight(name='tp', initializer='zeros')

    def update_state(self, y_true, y_pred, sample_weight=None):
        y_true = tf.cast(y_true, tf.bool)
        y_pred = tf.cast(y_pred, tf.bool)

        values = tf.logical_and(tf.equal(y_true, True), tf.equal(y_pred, True))
        values = tf.cast(values, self.dtype)
        if sample_weight is not None:
            sample_weight = tf.cast(sample_weight, self.dtype)
            values = tf.multiply(values, sample_weight)
        self.true_positives.assign_add(tf.reduce_sum(values))

    def result(self):
        return self.true_positives

    def reset_states(self):
        self.true_positives.assign(0)

m = BinaryTruePositives()
m.update_state([0, 1, 1, 1], [0, 1, 0, 0])
print('Intermediate result:', float(m.result()))

m.update_state([1, 1, 1, 1], [0, 1, 1, 0])
print('Final result:', float(m.result()))

```

The `add_metric()` API

When writing the forward pass of a custom layer or a subclassed model, you may sometimes want to log certain quantities on the fly, as metrics. In such cases, you can use the `add_metric()` method.

Let's say you want to log as metric the mean of the activations of a Dense-like custom layer. You could do the following:

```

class DenseLike(Layer):
    """y = w.x + b"""

    ...

    def call(self, inputs):
        output = tf.matmul(inputs, self.w) + self.b
        self.add_metric(tf.reduce_mean(output), aggregation='mean', name='activation_mean')
        return output

```

The quantity will then tracked under the name "activation_mean". The value tracked will be the average of the per-batch metric values (as specified by `aggregation='mean'`).

See [the `add_metric\(\)` documentation](https://keras.io/api/metrics/) for more details.