



tf.keras.metrics.AUC

See Stable See Nightly

 TensorFlow 1 version [./versions/r1.15/api_docs/python/tf/keras/metrics/AUC](https://www.tensorflow.org/api_docs/python/tf/keras/metrics/AUC)

 View source on GitHub <https://github.com/tensorflow/tensorflow/blob/v2.5.0/tensorflow/python/keras/metrics.py#L1940-L2410>

Approximates the AUC (Area under the curve) of the ROC or PR curves.

Inherits From: [Metric](https://www.tensorflow.org/api_docs/python/tf/keras/metrics/Metric) (https://www.tensorflow.org/api_docs/python/tf/keras/metrics/Metric), [Layer](https://www.tensorflow.org/api_docs/python/tf/keras/layers/Layer) (https://www.tensorflow.org/api_docs/python/tf/keras/layers/Layer), [Module](https://www.tensorflow.org/api_docs/python/tf/Module) (https://www.tensorflow.org/api_docs/python/tf/Module)

+ View aliases

Main aliases

[tf.metrics.AUC](https://www.tensorflow.org/api_docs/python/tf/keras/metrics/AUC) (https://www.tensorflow.org/api_docs/python/tf/keras/metrics/AUC)

Compat aliases for migration

See [Migration guide](https://www.tensorflow.org/guide/migrate) (<https://www.tensorflow.org/guide/migrate>) for more details.

[tf.compat.v1.keras.metrics.AUC](https://www.tensorflow.org/api_docs/python/tf/keras/metrics/AUC) (https://www.tensorflow.org/api_docs/python/tf/keras/metrics/AUC)

```
tf.keras.metrics.AUC(
    num_thresholds=200, curve='ROC',
    summation_method='interpolation', name=None, dtype=None,
    thresholds=None, multi_label=False, num_labels=None, label_weights=None,
    from_logits=False
)
```

Used in the notebooks

Used in the tutorials

- [Classification on imbalanced data](https://www.tensorflow.org/tutorials/structured_data/imbalanced_data) (https://www.tensorflow.org/tutorials/structured_data/imbalanced_data)
- [TF Lattice Premade Models](https://www.tensorflow.org/lattice/tutorials/premade_models) (https://www.tensorflow.org/lattice/tutorials/premade_models)
- [Client-efficient large-model federated learning via 'federated_select' and sparse aggregation](https://www.tensorflow.org/federated/tutorials/sparse_federated_learning) (https://www.tensorflow.org/federated/tutorials/sparse_federated_learning)

The AUC (Area under the curve) of the ROC (Receiver operating characteristic; default) or PR (Precision Recall) curves are quality measures of binary classifiers. Unlike the accuracy, and like cross-entropy losses, ROC-AUC and PR-AUC evaluate all the operational points of a model.

This classes approximates AUCs using a Riemann sum: During the metric accumulation phrase, predictions are accumulated within predefined buckets by value. The AUC is then computed by interpolating per-bucket averages. These buckets define the evaluated operational points.

This metric creates four local variables, `true_positives`, `true_negatives`, `false_positives` and `false_negatives` that are used to compute the AUC. To discretize the AUC curve, a linearly spaced set of thresholds is used to compute pairs of recall and precision values. The area under the ROC-curve is therefore computed using the height of the recall values by the false positive rate, while the area under the PR-curve is the computed using the height of the precision values by the recall.

This value is ultimately returned as `auc`, an idempotent operation that computes the area under a discretized curve of precision versus recall values (computed using the aforementioned variables). The `num_thresholds` variable controls the degree of discretization with larger numbers of thresholds more closely approximating the true AUC. The quality of the approximation may vary dramatically depending on `num_thresholds`. The `thresholds` parameter can be used to manually specify thresholds which split the predictions more evenly.

For a best approximation of the real AUC, `predictions` should be distributed approximately uniformly in the range `[0, 1]` ([/api_docs/python/tf/keras/metrics/if%20%60from_logits=False%60](https://www.tensorflow.org/api_docs/python/tf/keras/metrics/if%20%60from_logits=False%60)). The quality of the AUC approximation may be poor if this is not the case. Setting `summation_method` to 'minoring' or 'majoring' can help quantify the error in the approximation by providing lower or upper bound estimate of the AUC.

If `sample_weight` is `None`, weights default to 1. Use `sample_weight` of 0 to mask values.

Args	
num_thresholds	(Optional) Defaults to 200. The number of thresholds to use when discretizing the roc curve. Values must be > 1.
curve	(Optional) Specifies the name of the curve to be computed, 'ROC' [default] or 'PR' for the Precision-Recall-curve.
summation_method	(Optional) Specifies the Riemann summation method (https://en.wikipedia.org/wiki/Riemann_sum) used. 'interpolation' (default) applies mid-point summation scheme for ROC. For PR-AUC, interpolates (true/false) positives but not the ratio that is precision (see Davis & Goadrich 2006 for details); 'minoring' applies left summation for increasing intervals and right summation for decreasing intervals; 'majoring' does the opposite.
name	(Optional) string name of the metric instance.
dtype	(Optional) data type of the metric result.
thresholds	(Optional) A list of floating point values to use as the thresholds for discretizing the curve. If set, the <code>num_thresholds</code> parameter is ignored. Values should be in [0, 1]. Endpoint thresholds equal to <code>{-epsilon, 1+epsilon}</code> for a small positive epsilon value will be automatically included with these to correctly handle predictions equal to exactly 0 or 1.
multi_label	boolean indicating whether multilabel data should be treated as such, wherein AUC is computed separately for each label and then averaged across labels, or (when False) if the data should be flattened into a single label before AUC computation. In the latter case, when multilabel data is passed to AUC, each label-prediction pair is treated as an individual data point. Should be set to False for multi-class data.

num_labels	(Optional) The number of labels, used when multi_label' is True. If num_labels is not specified, then state variables get created on the first call to update_state. </td></tr><tr><td>label_weights</td><td> (Optional) list, array, or tensor of non-negative weights used to compute AUCs for multilabel data. When multi_label is True, the weights are applied to the individual label AUCs when they are averaged to produce the multi-label AUC. When it's False, they are used to weight the individual label predictions in computing the confusion matrix on the flattened data. Note that this is unlike class_weights in that class_weights weights the example depending on the value of its label, whereas label_weights depends only on the index of that label before flattening; therefore label_weights should not be used for multi-class data. </td></tr><tr><td>from_logits</td><td> boolean indicating whether the predictions (y_pred in update_state) are probabilities or sigmoid logits. As a rule of thumb, when using a keras loss, the from_logits constructor argument of the loss should match the AUC from_logits' constructor argument. <div data-bbox="142 365 277 387" data-label="Section-Header"><p>Standalone usage:</p></div> <div data-bbox="142 430 718 589" data-label="Code-Block"><pre>>>> m = tf.keras.metrics.AUC(num_thresholds=3) >>> m.update_state([0, 0, 1, 1], [0, 0.5, 0.3, 0.9]) >>> # threshold values are [0 - 1e-7, 0.5, 1 + 1e-7] >>> # tp = [2, 1, 0], fp = [2, 0, 0], fn = [0, 1, 2], tn = [0, 2, 2] >>> # recall = [1, 0.5, 0], fp_rate = [1, 0, 0] >>> # auc = (((1+0.5)/2)*(1-0))+ (((0.5+0)/2)*(0-0))) = 0.75 >>> m.result().numpy() 0.75</pre></div> <div data-bbox="142 651 582 750" data-label="Code-Block"><pre>>>> m.reset_state() >>> m.update_state([0, 0, 1, 1], [0, 0.5, 0.3, 0.9], ... sample_weight=[1, 0, 0, 1]) >>> m.result().numpy() 1.0</pre></div> <div data-bbox="142 790 339 815" data-label="Section-Header"><p>Usage with compile() API:</p></div> <div data-bbox="142 855 751 1034" data-label="Code-Block"><pre># Reports the AUC of a model outputting a probability. model.compile(optimizer='sgd', loss=tf.keras.losses.BinaryCrossentropy(), metrics=[tf.keras.metrics.AUC()]) # Reports the AUC of a model outputting a logit. model.compile(optimizer='sgd', loss=tf.keras.losses.BinaryCrossentropy(from_logits=True), metrics=[tf.keras.metrics.AUC(from_logits=True)])</pre></div> <div data-bbox="142 1086 207 1106" data-label="Section-Header"><p>Attributes</p></div> <div data-bbox="142 1128 611 1151" data-label="Table"><table><tr><td>thresholds</td><td>The thresholds used for evaluating AUC.</td></tr></table></div> <div data-bbox="142 1207 240 1236" data-label="Section-Header"><p>Methods</p></div> <div data-bbox="142 1263 339 1290" data-label="Section-Header"><p>interpolate_pr_auc</p></div> <div data-bbox="142 1305 904 1328" data-label="Text"><p>View source (https://github.com/tensorflow/tensorflow/blob/v2.5.0/tensorflow/python/keras/metrics.py#L2250-L2329)</p></div> <div data-bbox="142 1368 312 1391" data-label="Code-Block"><pre>interpolate_pr_auc()</pre></div> <div data-bbox="142 1431 635 1453" data-label="Text"><p>Interpolation formula inspired by section 4 of Davis & Goadrich 2006.</p></div> <div data-bbox="142 1471 782 1494" data-label="Text"><p>https://www.biostat.wisc.edu/~page/rocpr.pdf (https://www.biostat.wisc.edu/~page/rocpr.pdf)</p></div> <div data-bbox="142 1512 708 1534" data-label="Text"><p>Note here we derive & use a closed formula not present in the paper as follows:</p></div> <div data-bbox="142 1550 391 1572" data-label="Equation-Block"><p>Precision = TP / (TP + FP) = TP / P</p></div> <div data-bbox="142 1590 1441 1612" data-label="Text"><p>Modeling all of TP (true positive), FP (false positive) and their sum P = TP + FP (predicted positive) as varying linearly within each interval [A, B] between successive thresholds, we get</p></div> <div data-bbox="142 1628 1008 1653" data-label="Equation-Block"><p>Precision slope = dTP / dP = (TP_B - TP_A) / (P_B - P_A) = (TP - TP_A) / (P - P_A) Precision = (TP_A + slope * (P - P_A)) / P</p></div> <div data-bbox="142 1668 587 1693" data-label="Text"><p>The area within the interval is (slope / total_pos_weight) times</p></div> <div data-bbox="142 1709 1067 1731" data-label="Equation-Block"><p>int_A^B(Precision.dP) = int_A^B((TP_A + slope * (P - P_A)) * dP / P) int_A^B(Precision.dP) = int_A^B(slope * dP + intercept * dP / P)</p></div> <div data-bbox="142 1749 638 1771" data-label="Text"><p>where intercept = TP_A - slope * P_A = TP_B - slope * P_B, resulting in</p></div> <div data-bbox="142 1789 600 1812" data-label="Equation-Block"><p>int_A^B(Precision.dP) = TP_B - TP_A + intercept * log(P_B / P_A)</p></div> <div data-bbox="142 1830 665 1852" data-label="Text"><p>Bringing back the factor (slope / total_pos_weight) we'd put aside, we get</p></div> <div data-bbox="142 1870 571 1892" data-label="Equation-Block"><p>slope * [dTP + intercept * log(P_B / P_A)] / total_pos_weight</p></div> <div data-bbox="142 1910 336 1933" data-label="Equation-Block"><p>where dTP == TP_B - TP_A.</p></div> <div data-bbox="142 1948 580 1971" data-label="Text"><p>Note that when P_A == 0 the above calculation simplifies into</p></div> <div data-bbox="142 1989 638 2011" data-label="Equation-Block"><p>int_A^B(Precision.dTP) = int_A^B(slope * dTP) = slope * (TP_B - TP_A)</p></div> <div data-bbox="142 2029 914 2051" data-label="Text"><p>which is really equivalent to imputing constant precision throughout the first bucket having >0 true positives.</p></div> <div data-bbox="71 2172 1524 2201" data-label="Page-Footer"><div>https://www.tensorflow.org/api_docs/python/tf/keras/metrics/AUC</div><div>2/3</div></div>	thresholds	The thresholds used for evaluating AUC.
thresholds	The thresholds used for evaluating AUC.		

Returns	
pr_auc	an approximation of the area under the P-R curve.

reset_state

[View source](https://github.com/tensorflow/tensorflow/blob/v2.5.0/tensorflow/python/keras/metrics.py#L2383-L2391) (https://github.com/tensorflow/tensorflow/blob/v2.5.0/tensorflow/python/keras/metrics.py#L2383-L2391)

```
reset_state()
```

Resets all of the metric state variables.

This function is called between epochs/steps, when a metric is evaluated during training.

result

[View source](https://github.com/tensorflow/tensorflow/blob/v2.5.0/tensorflow/python/keras/metrics.py#L2331-L2381) (https://github.com/tensorflow/tensorflow/blob/v2.5.0/tensorflow/python/keras/metrics.py#L2331-L2381)

```
result()
```

Computes and returns the metric value tensor.

Result computation is an idempotent operation that simply calculates the metric value using the state variables.

update_state

[View source](https://github.com/tensorflow/tensorflow/blob/v2.5.0/tensorflow/python/keras/metrics.py#L2186-L2248) (https://github.com/tensorflow/tensorflow/blob/v2.5.0/tensorflow/python/keras/metrics.py#L2186-L2248)

```
update_state(
    y_true, y_pred, sample_weight=None
)
```

Accumulates confusion matrix statistics.

Args	
y_true	The ground truth values.
y_pred	The predicted values.
sample_weight	Optional weighting of each example. Defaults to 1. Can be a <code>Tensor</code> whose rank is either 0, or the same rank as <code>y_true</code> , and must be broadcastable to <code>y_true</code> .
Returns	
Update op.	

Except as otherwise noted, the content of this page is licensed under the [Creative Commons Attribution 4.0 License](https://creativecommons.org/licenses/by/4.0/) (https://creativecommons.org/licenses/by/4.0/), and code samples are licensed under the [Apache 2.0 License](https://www.apache.org/licenses/LICENSE-2.0) (https://www.apache.org/licenses/LICENSE-2.0). For details, see the [Google Developers Site Policies](https://developers.google.com/site-policies) (https://developers.google.com/site-policies). Java is a registered trademark of Oracle and/or its affiliates. Some content is licensed under the [numpy license](https://numpy.org/doc/stable/license.html) (https://numpy.org/doc/stable/license.html).

Last updated 2021-06-24 UTC.

